

AI Bot Access-Log Analysis

Turn raw server access logs into a defensible AI-visibility metric — by separating training crawls from search crawls from bots acting on a live user's behalf.

Source: **Chris Green**, Technical Director & Senior Consultant, Torque Partnership — Unscripted SEO Podcast.

OBJECTIVE

Produce a repeatable measure of AI visibility from data you already own. Count AI bot hits per content segment, classify each hit by what the bot was doing — training crawl, search crawl, or a fetch triggered by a live human asking an assistant — and track that third category over time. The third number is the deliverable. It is not revenue and must never be presented as revenue. It is the closest thing to a visibility signal you can generate from your own server, without buying a prompt-tracking tool and without arguing about whose share-of-voice methodology is right.

"one of the other treasure troves of information is access logs." — Chris Green

KEY STEPS

1 Settle the data-access question before you promise the deliverable.

Find out what you may do with the logs, and specifically whether you may put them in front of an AI tool. This kills projects after they are sold. Chris: *"The thing I'm reluctant to do at the moment, and that often depends on my MSAs and my security arrangements and stuff, is giving direct access to your data."* Get it in writing, then design around the tier you got.

2 Segment your content before you read a single log line.

Group URLs into meaningful segments — templates, categories, docs tree, money pages. The segment, not the URL, is the unit of analysis: *"take an average of like bot hits to a segment of pages"*. One regex per segment; write it down and reuse the same one forever.

3 Baseline the average bot hits per segment.

Pull a fixed window and record average hits per segment, per bot, per day. The boring step everyone skips, and the one that makes every later number mean something. *"That's probably one of the most simple things you can do."*

4 Run it as a split test, not a before-and-after.

"Say you run like a split test. So you say I've updated some pages and not others segment those pages. Take a grab of your log files at a certain date and then compare it." The untouched segment is your control against seasonality and against a bot operator simply changing its own crawl budget.

5 Classify every hit by what the bot was doing, not just who sent it.

Three buckets: training crawl, search crawl, and a bot acting on a live user's behalf. *"So you can at least understand is this bot just crawling for training data or is this bot acting on the user's behalf?"* OpenAI, Perplexity and Anthropic publish *"fairly good detailed levels"* of documentation; *"Google are the least helpful with how they designate AI bots"*. Keep the mapping table in version control.

6 Verify the user agent against published IP ranges before you trust anything.

A user-agent string is a claim, not a fact. *"Is it a valid chat GPT hit or is it someone that's spoofing it? And without telling Claude to check the IP range and the other network conditions, you may start analyzing bot activity on the wrong bot. It could just be a spoofer."* Make verification a mandatory filter step before aggregation, not a spot check after.

7 Move to a queryable store once the volume gets real.

"Log files can be enormous." Chris's options: *"you need to use like a log file analysis tool or a tool like Splunk or you need to push it into a database and use BigQuery or SQL or to query it."* His preferred route is BigQuery with reporting on top — *"You can kind of use Data Studio as like a front end"*. Then the AI assist is a SQL-writing job, not a data-access job.

8 Report user-triggered hits as a correlation, and bound the claim out loud.

"acting on the user's behalf is a proxy for visibility or being surfaced at least by an AI tool." Then, immediately: *"that's not attached to revenue, but you can say, well, this correlates with our involvement. We are making your site more findable potentially. And that's a good thing to ground yourself on."* Volunteering the limit is what makes the number survive scrutiny.

CAUTIONARY NOTES

- **An unverified user agent points you at the wrong bot entirely.** This invalidates the report rather than just adding noise — which is why IP verification is a step, not a tip. *"It could just be a spoofer."*
- **The model will find a trend that is not there.** *"It's very, very easy to start hallucinating things or it might start seeing a trend in the data that's not really there."* Ask for the query and the raw counts; sanity-check the counts before anything reaches a slide.
- **Most of a log file is noise.** *"depending on your log format, you can have a lot of data in there that's just noise."* Decide your exclusions — static assets, health checks, your own monitoring — and write them into the query so they apply identically every run.
- **The bot taxonomy moves under you.** On how many bot types OpenAI runs: *"I think there's three now. I think there might even be more. There's user search and then training for chat gpt. I say that with confidence it changes all the time."* Re-read the vendors' published bot lists on a schedule.
- **Google will not make this easy.** *"Google are the least helpful with how they designate AI bots"*. Expect a coverage gap on the largest player and label it as a gap rather than papering over it.
- **Never attach this number to revenue.** It is a proxy and a leading indicator. Presented as attributable performance, it becomes exactly the argument-generating metric this play exists to avoid — and you lose that argument the first time it dips.
- **BigQuery is not free.** *"Obviously it costs, you have to pay data processing."* Partition by date, select only the columns you need, and price the recurring query before you promise a monthly report.
- **Client security terms may forbid the AI-assisted version.** Some environments allow the analysis but not the tooling. Know which one you are in before kickoff, not after.

TIPS FOR EFFICIENCY

- **Know how you would do the analysis without the AI first.** *"if you kind of know how you would do it without Claude, you can get Claude to do it much more effectively for you"*. If you cannot describe the query in words, you cannot review the one you get back.
- **Give it your field names, then ask for the query — not the answer.** Chris's prompt shape: *"segment my pages on this regex and then I want to understand what my running average for bot hits are. I want to find who my outliers are. Like help me write a SQL query that can help identify that"*. This also keeps the raw data out of the tool when your MSA requires that.
- **Ask for outliers, not only averages.** The running average gives you the trend; the outliers tell you which segment to go look at. The outlier list is a free ranked to-do list.
- **Store the segmentation regex, the exclusion rules and the query as one saved artifact.** You re-run the same thing every month. If any of the three drifts, the comparison is worthless and nothing warns you.
- **Put a dashboard on it once, not a deck every month.** Looker Studio (Data Studio) over BigQuery turns a recurring analysis into a link you send. The build cost is one-off.
- **Pair the chart with the sentence.** The number does not sell itself — the framing does. Keep *"acting on the user's behalf is a proxy for visibility"* and *"this correlates with our involvement"* on the same slide as the trend line, every time.

Source — Chris Green, Torque Partnership. Technical Director and Senior Consultant, running global SEO, GEO, measurement and AI search programmes for Fortune 500 brands. He framed access logs as the measurement layer almost nobody is selling — *"in my experience, very few people are bringing that kind of value"* — and prefers user-triggered AI bot hits over prompt-tracking share of voice because *"If all of the other metrics you have will be share of voice or visibility and they'll be based on prompt tracking and then it gets horribly complex and easy to argue about."* Every quote here is his, verbatim from the episode; the caveats are his too.

Full interview: unscriptedseo.com/chris-green-do-googles-patents-still-hold-up-in-an-agentic-web/

Read the SOP online: seorcade.com/seo-sop-ai-bot-access-log-analysis/

More episodes: unscriptedseo.com — 150+ interviews with SEO practitioners.