

Publish a Machine-Readable Companion File for AI Crawlers

Bots now do most of the crawling, and they have a preferred content type — so give them one. Publish a stripped plain-text/markdown version of your highest-stakes pages, link to it, and watch your logs. Distilled from Dan Kurtz's Unscripted SEO interview with host Jeremy Rivera.

Objective

Make the facts you care most about — what you do, what you charge, who you serve — cheap and unambiguous for an AI crawler to read, by publishing a plain-text/markdown companion version of your highest-stakes pages at stable URLs and linking to them from the canonical HTML. Not a replacement for your site, not a ranking hack: one more format in the multi-purposing rotation you already run, aimed at the audience doing most of the crawling. The downside is bounded — as Dan puts it, “the worst thing that can happen is they pass it. They fly over it.”

Key Steps

1. **List the pages whose facts you most need repeated back correctly.** Not the whole site — the handful where a wrong answer costs money: pricing, services, coverage area, who you serve, key specs. Five to ten URLs is a real first pass.
2. **Produce a stripped, structured text version of each.** Clear headings, bullet lists, plain declarative sentences. No design, nav, modals or cookie banner. A full HTML page costs a bot roughly 3× longer to load than markdown costs in tokens and time — and markdown is already organized the way the machine likes it. Every fact identical to the canonical page.
3. **Publish each file at a stable URL, like any other asset.** One convention, never moved — simplest is canonical path plus `.md` (`/pricing/` → `/pricing.md`). Serve as plain text, un-redirected, not blocked in robots.txt. (*inferred — Dan specified the files, not a naming scheme.*)
4. **Link the file from the canonical page so crawlers discover it.** The step people skip, and the whole mechanism: “that’s just a matter of directing them to it next time they come to the website and saying, hey, this file exists.” A visible footer link is enough; add them to the sitemap plus an index file for a second discovery path (*inferred*).
5. **Baseline your bot traffic before anything changes.** Pull 30–90 days of server logs and split requests by user agent into traditional crawlers (Googlebot, Bingbot) vs AI crawlers/retrievers (GPTBot, ClaudeBot, PerplexityBot, Google-Extended). Write the ratio down — without a “before,” Step 6 tells you nothing.
6. **Verify pickup by watching the log ratio, not rankings.** Weeks later, re-run the same split: are agents requesting the `.md` URLs, is the AI-vs-traditional ratio moving, and which files get hit most. Dan does exactly this — “it’s just looking at server logs... got my own sites that I can see those kinds of tests on.”
7. **Put the companion file in your publishing checklist.** When the canonical page changes, the companion changes in the same ticket. A stale machine-readable file is a confidently wrong answer.
8. **Expand deliberately, on evidence.** Add the next tier only once the logs show the first tier being read — a format for a confirmed audience, not files generated at scale for their own sake.

Cautionary Notes

- **Companion, never replacement.** Google’s public guidance is not to serve markdown in place of your pages. The HTML page stays canonical and intact.
- **Don’t drift into cloaking.** Same facts, claims and prices as the page it mirrors. The moment they diverge you’ve built a cloaking problem.
- **Accept it might do nothing.** Budget an afternoon, not a quarter.
- **Don’t generalize one practitioner’s logs.** Dan’s ~60% down (Googlebot) / ~80% up (AI bots) figures are *his own server logs*. The verified industry number is Cloudflare’s majority-bot crossover (Matthew Prince, 3 June 2026, ~57.5% automated).
- **Stale beats nothing only until it’s wrong.** An outdated price in a file built to be quoted verbatim is worse than no file.

Tips for Efficiency

- Convert canonical HTML to markdown programmatically, then edit down — don’t rewrite by hand.
- Pick the naming/link convention before the first file; retrofitting across thirty files is where this dies.
- Save the log query, not the result, so before/after comparisons are genuinely identical.
- Sort the AI-bot bucket by file, not just count — it’s a free ranked to-do list of what to convert next.
- Add it to the multi-purposing checklist you already run (infographic, short, press release...) so it costs nothing organizationally.
- Periodically ask the major assistants what your company charges and who it serves. Better answers are the real outcome; crawl ratio is only a proxy.

Sidebar — the SSL re-crawl trick is NOT part of this SOP. Dan mentioned that reissuing your SSL certificate appears to push AI bots to re-crawl, and flagged its limits in the same breath: “I don’t know what the cadence is, I don’t know what the frequency is,” and someone will “spam it to death and find out that their website doesn’t take bots anymore.” His label: “*this is a table saw — watch your fingers and make sure you can still count to ten.*” Listed as an optional, unverified accelerator only. Nothing above depends on it.

Source: Dan Kurtz, founder of Miss Pepper AI — who reframed serving machine-readable files as ordinary multi-purposing rather than a new tactic: “Bots have a preferred content type. Technically, they’re just an additional customer.”

Full interview: unscriptedseo.com/dan-kurtz-bots-as-a-customer-segment-markdown-agentic-flywheel/ · More: unscriptedseo.com