

Mine Your 404 Report for AI-Hallucinated URLs

From the Unscripted SEO interview with **Stephanie Long** — Mrs. SEO / Junction Jack · host Jeremy Rivera · "SEO Is the Batter to Your AI"

OBJECTIVE

Split your 404s into links that used to work and URLs that never existed, because an AI assistant invented them. The second pile is a demand signal: each invented URL is a page a model decided you should have, handed to a real person, who clicked through to nothing. You end with a ranked build list of pages that already have traffic aimed at them, plus redirects that stop the near-misses leaking. One afternoon for the first pass, an hour a month after.

KEY STEPS

1. **Open Google Search Console → Pages → Not indexed → "Not found (404)."** Already deduplicated and counted. Pull server logs and any 404 plugin too, since they catch hits GSC never sees. Window: last 90 days.
2. **Sort by hits and export.** Volume is the ranking signal. Invented once is noise; invented repeatedly is recurring demand. Work top down, stop when counts go flat.
3. **Split into "used to exist" and "never existed."** Check the CMS, sitemap history, and the Wayback Machine. Anything never published is a candidate invented URL. (*inferred*)
4. **Fix the "used to exist" pile the ordinary way and set it aside.** Redirect, restore, or leave dead. Routine housekeeping — getting it out of the list stops it obscuring the interesting half.
5. **Read the invented URLs as a topic list.** Group the slugs. The model picked a path that fits your structure and a topic it believed you would cover. Repeated shapes tell you what buyers are being told you offer. (*inferred*)
6. **Build the invented page at that exact URL — where it is one you would actually want.** The path matters as much as the content, because the traffic is already addressed to it. If it misrepresents what you do, do not manufacture it just because a model asked.
7. **Redirect the near-misses to the closest real page.** Many invented slugs are variants of something live. Those get a 301, not a new page. Build only when there is nothing genuine to point at.
8. **Re-pull monthly and keep score.** Invented URLs shift as models retrain. Track hits before and after on the pages you build. (*inferred*)

CAUTIONARY NOTES

- **Do not build every invented URL.** The filter is "would I want this page anyway," not "did a machine ask for it."
- **A 301 beats a thin new page.** Building a second page for the same intent splits your own signals.
- **GSC does not see everything.** People arriving from a chat answer often never touch Google, so logs show hits the report misses.
- **Counts are small by SEO standards.** This is high intent, not scale: they were told you had the answer and hit a wall.
- **Some invented links are not on your domain.** Stephanie clicked every citation on one AI query and found all of them broken. You cannot fix those from your own 404 report.
- **It is not an AI ranking win.** You are catching traffic already being sent. Whether the model then cites the real page is a separate measurement.

TIPS FOR EFFICIENCY

- Filter logs by AI assistant referrers and user agents first — a much cleaner starting list, one grep.
- Ask the assistants directly. Prompt for a resource you would plausibly publish and see which URL comes back, before anyone hits it.
- Keep the built pages in one tracked analytics group so before-and-after is a single screenshot.
- Run the same pass on a competitor's domain during a pitch: invented URLs there are category pages nobody has built yet.
- Batch it with normal broken-link cleanup. Two piles instead of one costs almost nothing extra.

Source: The play is Jeremy Rivera's, citing an Ahrefs post by Tim Soulo or Patrick Stox. **Stephanie Long** (Mrs. SEO / Junction Jack) endorsed it on tape from the other side.

"Check your 404s on your site. Not because it's broken, but because it was imagineered by ChatGPT... you might as well add that resource, to plug in the traffic that they're going to send to imaginary land."